

Assignment 9.2

Problem 1:

Create an HBase table named 'clicks' with a column family 'hits' such that it should be able to store last 5 values of qualifiers inside 'hits' column family.

```
create 'clicks','hits'
```

```
put 'clicks','192.168.92.1','hits:values','siteA'
```

```
put 'clicks','192.168.92.2','hits:values','siteB'
```

```
put 'clicks','192.168.92.3','hits:values','siteC'
```

```
put 'clicks','192.168.92.4','hits:values','siteD'
```

```
put 'clicks','192.168.92.5','hits:values','siteE'
```

```
scan 'clicks'
```

```
hbase(main):002:0> create 'clicks','hits'
0 row(s) in 2.0120 seconds

=> Hbase::Table - clicks
hbase(main):003:0> put 'clicks','192.168.92.1','hits:values','siteA'
0 row(s) in 0.3850 seconds

hbase(main):004:0> put 'clicks','192.168.92.2','hits:values','siteB'
0 row(s) in 0.2220 seconds

hbase(main):005:0> put 'clicks','192.168.92.3','hits:values','siteC'
0 row(s) in 0.0300 seconds

hbase(main):006:0> put 'clicks','192.168.92.4','hits:values','siteD'
0 row(s) in 0.0120 seconds

hbase(main):007:0> put 'clicks','192.168.92.5','hits:values','siteE'
0 row(s) in 0.0170 seconds

hbase(main):008:0> scan 'clicks'
ROW                                COLUMN+CELL
192.168.92.1                      column=hits:values, timestamp=1513556766000, value=siteA
192.168.92.2                      column=hits:values, timestamp=1513556766294, value=siteB
192.168.92.3                      column=hits:values, timestamp=1513556766502, value=siteC
192.168.92.4                      column=hits:values, timestamp=1513556766590, value=siteD
192.168.92.5                      column=hits:values, timestamp=1513556766639, value=siteE
5 row(s) in 0.1120 seconds

hbase(main):009:0> █
```

Problem 2:

Add few records in the table and update some of them. Use IP Address as row-key. Scan the table to view if all the previous versions are getting displayed

```
alter 'clicks', NAME=>'hits', VERSIONS=>2
```

```
put 'clicks','192.168.92.1','hits:values','siteF'
```

```
put 'clicks','192.168.92.1','hits:values','siteG'
```

```
scan 'clicks'
```

```
scan 'clicks',{VERSIONS=>2}
```

```
hbase(main):009:0> alter 'clicks', NAME=>'hits', VERSIONS=>2
Updating all regions with the new schema...
0/1 regions updated.
1/1 regions updated.
Done.
0 row(s) in 2.6370 seconds

hbase(main):010:0> put 'clicks','192.168.92.1','hits:values','siteF'
0 row(s) in 0.0780 seconds

hbase(main):011:0> put 'clicks','192.168.92.1','hits:values','siteG'
0 row(s) in 0.0520 seconds

hbase(main):012:0> scan 'clicks'
ROW                                COLUMN+CELL
192.168.92.1                       column=hits:values, timestamp=1513560530434, value=siteG
192.168.92.2                       column=hits:values, timestamp=1513556766294, value=siteB
192.168.92.3                       column=hits:values, timestamp=1513556766502, value=siteC
192.168.92.4                       column=hits:values, timestamp=1513556766590, value=siteD
192.168.92.5                       column=hits:values, timestamp=1513556769639, value=siteE
5 row(s) in 0.0680 seconds

hbase(main):013:0> scan 'clicks',{VERSIONS=>1}
ROW                                COLUMN+CELL
192.168.92.1                       column=hits:values, timestamp=1513560530434, value=siteG
192.168.92.2                       column=hits:values, timestamp=1513556766294, value=siteB
192.168.92.3                       column=hits:values, timestamp=1513556766502, value=siteC
192.168.92.4                       column=hits:values, timestamp=1513556766590, value=siteD
192.168.92.5                       column=hits:values, timestamp=1513556769639, value=siteE
5 row(s) in 0.1020 seconds

hbase(main):014:0> scan 'clicks',{VERSIONS=>2}
ROW                                COLUMN+CELL
192.168.92.1                       column=hits:values, timestamp=1513560530434, value=siteG
192.168.92.1                       column=hits:values, timestamp=1513560526482, value=siteF
192.168.92.2                       column=hits:values, timestamp=1513556766294, value=siteB
192.168.92.3                       column=hits:values, timestamp=1513556766502, value=siteC
192.168.92.4                       column=hits:values, timestamp=1513556766590, value=siteD
```

```
alter 'clicks', NAME=>'hits', VERSIONS=>3
```

```
put 'clicks','192.168.92.1','hits:values','siteW'
```

```
put 'clicks','192.168.92.2','hits:values','siteX'
```

```
put 'clicks','192.168.92.3','hits:values','siteY'
```

```
put 'clicks','192.168.92.4','hits:values','siteZ'
```

```
scan 'clicks',{VERSIONS=>4}
```

```

hbase(main):016:0> alter 'clicks', NAME=>'hits', VERSIONS=>3
Updating all regions with the new schema...
0/1 regions updated.
1/1 regions updated.
Done.
0 row(s) in 2.2460 seconds

hbase(main):017:0> put 'clicks','192.168.92.1','hits:values','siteW'
0 row(s) in 0.0360 seconds

hbase(main):018:0> put 'clicks','192.168.92.2','hits:values','siteX'
0 row(s) in 0.0080 seconds

hbase(main):019:0> put 'clicks','192.168.92.3','hits:values','siteY'
0 row(s) in 0.0050 seconds

hbase(main):020:0> put 'clicks','192.168.92.4','hits:values','siteZ'
0 row(s) in 0.0120 seconds

hbase(main):021:0> scan 'clicks', {VERSIONS => 4}
ROW                                COLUMN+CELL
192.168.92.1                       column=hits:values, timestamp=1513560803430, value=siteW
192.168.92.1                       column=hits:values, timestamp=1513560530434, value=siteG
192.168.92.1                       column=hits:values, timestamp=1513560526482, value=siteF
192.168.92.2                       column=hits:values, timestamp=1513560803470, value=siteX
192.168.92.2                       column=hits:values, timestamp=1513556766294, value=siteB
192.168.92.3                       column=hits:values, timestamp=1513560803528, value=siteY
192.168.92.3                       column=hits:values, timestamp=1513556766502, value=siteC
192.168.92.4                       column=hits:values, timestamp=1513560805407, value=siteZ
192.168.92.4                       column=hits:values, timestamp=1513556766590, value=siteD
192.168.92.5                       column=hits:values, timestamp=1513556769639, value=siteE
5 row(s) in 0.1030 seconds

hbase(main):022:0> █

```